

Uncertainty Estimation for High-dimensional Nonparametric Forecasting

Nuwani Palihawadana

Joint work with : Rob Hyndman, Xiaoqian Wang

1 July 2025

Nonparametric model

$$y_t = f(\mathbf{a}_t, \mathbf{a}_{t-1}, \dots, \mathbf{a}_{t-s}, y_{t-1}, \dots, y_{t-k}) + \varepsilon_t,$$

- y_t – observed response at time t
- f – arbitrary function
- $\mathbf{a}_t$ – vector of exogenous variables at time t
- ε_t – stationary error with mean zero and constant variance σ^2

We also define:

- $\mathbf{z}_t = (y_t, \mathbf{x}_t)$ – observation at time t
- $\mathbf{x}_t = (\mathbf{a}_t, \mathbf{a}_{t-1}, \dots, \mathbf{a}_{t-s}, y_{t-1}, \dots, y_{t-k})$ – vector of all predictors at time t .
- $\hat{y}_t = \hat{f}(\mathbf{x}_t)$ – estimated model
- $e_t = y_t - \hat{y}_t$ – model residuals

Sparse Multiple Index (SMI) model

$$y_i = \beta_0 + \sum_{j=1}^p g_j(\boldsymbol{\alpha}_j^T \mathbf{x}_{ij}) + \sum_{k=1}^d f_k(w_{ik}) + \boldsymbol{\theta}^T \mathbf{u}_i + \varepsilon_i, \quad i = 1, \dots, n,$$

- y_i – univariate response
- $\mathbf{x}_{ij} \in \mathbb{R}^{\ell_j}$, $j = 1, \dots, p$ – p subsets of predictors entering indices
- $\boldsymbol{\alpha}_j$ – ℓ_j -dimensional vectors of index coefficients
- g_j, f_k – smooth nonlinear functions
- Additional predictors :
 - ▶ w_{ik} – nonlinear
 - ▶ $\mathbf{u}_i$ – linear

Sparse Multiple Index (SMI) model

$$y_i = \beta_0 + \sum_{j=1}^p g_j(\boldsymbol{\alpha}_j^T \mathbf{x}_{ij}) + \sum_{k=1}^d f_k(w_{ik}) + \boldsymbol{\theta}^T \mathbf{u}_i + \varepsilon_i, \quad i = 1, \dots, n,$$

- y_i – univariate response
- $\mathbf{x}_{ij} \in \mathbb{R}^{\ell_j}$, $j = 1, \dots, p$ – p subsets of predictors entering indices
- $\boldsymbol{\alpha}_j$ – ℓ_j -dimensional vectors of index coefficients
- g_j, f_k – smooth nonlinear functions
- Additional predictors :
 - ▶ w_{ik} – nonlinear
 - ▶ $\mathbf{u}_i$ – linear

Allow elements equal to zero in $\boldsymbol{\alpha}_j$ – "Sparse"

Sparse Multiple Index (SMI) model

$$y_i = \beta_0 + \sum_{j=1}^p g_j(\boldsymbol{\alpha}_j^T \mathbf{x}_{ij}) + \sum_{k=1}^d f_k(w_{ik}) + \boldsymbol{\theta}^T \mathbf{u}_i + \varepsilon_i, \quad i = 1, \dots, n,$$

- y_i – univariate response
- $\mathbf{x}_{ij} \in \mathbb{R}^{\ell_j}$, $j = 1, \dots, p$ – p subsets of predictors entering indices
- $\boldsymbol{\alpha}_j$ – ℓ_j -dimensional vectors of index coefficients
- g_j, f_k – smooth nonlinear functions
- Additional predictors :
 - ▶ w_{ik} – nonlinear
 - ▶ $\mathbf{u}_i$ – linear

Both " p " and the predictor grouping among indices are unknown.

Overlapping of predictors among indices is not allowed.

- Nonparametric additive model with backward elimination (Backward):
 - ▶ No linear combinations (indices)
 - ▶ Fully additive

- Nonparametric additive model with backward elimination (Backward):
 - ▶ No linear combinations (indices)
 - ▶ Fully additive
- Groupwise Additive Index Model (GAIM):
 - ▶ Predefined predictor groups
 - ▶ No overlapping predictors among groups

- Nonparametric additive model with backward elimination (Backward):
 - ▶ No linear combinations (indices)
 - ▶ Fully additive
- Groupwise Additive Index Model (GAIM):
 - ▶ Predefined predictor groups
 - ▶ No overlapping predictors among groups
- Projection Pursuit Regression model (PPR):
 - ▶ All predictors enter all indices

Forecast uncertainty

- Uncertainty of a forecast → **Prediction Interval (PI)**

Forecast uncertainty

- Uncertainty of a forecast → **Prediction Interval (PI)**
- Theoretical $100(1 - \alpha)\%$ prediction interval:

$$\hat{y}_{t+h|t} \pm z_{\alpha/2} \times \hat{\sigma}_h,$$

where

- ▶ y – time series $y_1, \dots, y_T$
- ▶ $\hat{y}_{t+h|t}$ – h -step-ahead point forecast for y_{t+h} given observations up to t
- ▶ $z_{\alpha/2}$ – $\alpha/2$ quantile of standard normal distribution
- ▶ $\hat{\sigma}_h$ – estimate of std. deviation of h -step forecast distribution

Forecast uncertainty

- Uncertainty of a forecast → **Prediction Interval (PI)**
- Theoretical $100(1 - \alpha)\%$ prediction interval:

$$\hat{y}_{t+h|t} \pm z_{\alpha/2} \times \hat{\sigma}_h,$$


where

- ▶ y – time series $y_1, \dots, y_T$
 - ▶ $\hat{y}_{t+h|t}$ – h -step-ahead point forecast for y_{t+h} given observations up to t
 - ▶ $z_{\alpha/2}$ – $\alpha/2$ quantile of standard normal distribution
 - ▶ $\hat{\sigma}_h$ – estimate of std. deviation of h -step forecast distribution
- Main issue:
 - ▶ Difficult to analytically calculate h -step forecast variances for $h > 1$

Block bootstrapping

- Randomly resample blocks from the historical model residuals, and join together
- Retains serial correlation in the data

BB with time series cross-validation



- tr : length of training window
- H : forecast horizon

BB with time series cross-validation

Step 1: Split the data into an initial training window, $\{\mathbf{z}_1, \dots, \mathbf{z}_{tr}\}$, and an initial test window, $\{\mathbf{z}_{tr+1}, \dots, \mathbf{z}_{tr+H}\}$.

Step 2: Train forecasting model on training window.

Step 3: Obtain the series of in-sample residuals, $e_1, \dots, e_{tr}$.

Step 4: Perform BB by using residuals series; generate several bootstrapped series.

Step 5: Obtain H -step-ahead simulated future values, $\hat{y}_{tr+1|tr}, \dots, \hat{y}_{tr+H|tr}$, for each bootstrapped series.

Step 6: Calculate quantiles of the sets of simulated future values at each horizon.

Step 7: Iteratively roll training window forward by one observation; repeat steps 2–6 until prediction intervals have been constructed for the entire test set.

Conformal prediction

- A distribution-free approach (Vovk et al. 2005)
- Relies only on the assumption of **exchangeability of data**
- Provides theoretical coverage guarantees

Conformal prediction

- A distribution-free approach (Vovk et al. 2005)
- Relies only on the assumption of **exchangeability of data**
- Provides theoretical coverage guarantees

Split Conformal Prediction (SCP) :

- A holdout method for generating prediction intervals
 - ▶ **Training set** – train forecasting model
 - ▶ **Calibration set** – calculate forecast errors (*nonconformity scores*)
 - ▶ **Test set** – obtain prediction intervals

Weighted Split Conformal Prediction (WSCP) (Barber et al. 2023) :

- Weighting quantiles using fixed (data-independent) weights

CP methods for non-exchangeable data


Weighted Split Conformal Prediction (WSCP) (Barber et al. 2023) :

- Weighting quantiles using fixed (data-independent) weights

Adaptive Conformal Prediction (ACP) (Gibbs & Candès 2021) :

- Update nominal α based on achieved coverage

CP with time series cross-validation (Wang & Hyndman 2024)



- tr : length of training window
- tc : length of calibration window
- H : forecast horizon

CP with time series cross-validation (Wang & Hyndman 2024)

Step 1: Split data into an initial training window, $\{\mathbf{z}_1, \dots, \mathbf{z}_{tr}\}$, an initial calibration window, $\{\mathbf{z}_{tr+1}, \dots, \mathbf{z}_{tr+tc}\}$, and an initial test window, $\{\mathbf{z}_{tr+tc+1}, \dots, \mathbf{z}_{tr+tc+H}\}$.

Step 2: Estimate forecasting model on initial training window. Obtain H -step-ahead forecasts, $\hat{y}_{tr+h|tr}$, and forecast errors $s_{tr+h|tr} = y_{tr+h} - \hat{y}_{tr+h|tr}$, for $h = 1, \dots, H$.

Step 3: Perform cross-validation forecasting, while repeating step 2, until forecast errors are obtained for entire initial calibration window.


Step 4: Apply the conformal prediction method of interest: SCP, WSCP, or ACP, to compute H -step prediction intervals.

Step 5: Iteratively roll training window and calibration window forward by one observation; repeat steps 2–4 until prediction intervals are obtained for the entire test set.

Conformal bootstrap

- **A natural integration of BB and SCP**
- **Exploits the strengths of both the methods**
 - ▶ Preserves temporal dependencies
 - ▶ Accounts for additional uncertainty brought into the process by ex-ante forecasting

CB with time series cross-validation



- tr : length of training window
- tc : length of calibration window
- H : forecast horizon

CB with time series cross-validation

Step 1: Split data into an initial training window, $\{\mathbf{z}_1, \dots, \mathbf{z}_{tr}\}$, an initial calibration window, $\{\mathbf{z}_{tr+1}, \dots, \mathbf{z}_{tr+tc}\}$, and an initial test window, $\{\mathbf{z}_{tr+tc+1}, \dots, \mathbf{z}_{tr+tc+H}\}$.

Step 2: Estimate forecasting model on initial training window. Obtain H -step-ahead forecasts, $\hat{y}_{tr+h|tr}$, and forecast errors $s_{tr+h|tr} = y_{tr+h} - \hat{y}_{tr+h|tr}$, for $h = 1, \dots, H$.

Step 3: Perform cross-validation forecasting, while repeating step 2, until forecast errors are obtained for entire initial calibration window.


Step 4: Use the model estimated on the most recent training window to generate H -step-ahead simulated future values.

- Autoregressive: block bootstrap 1-step forecast errors; block size = H .
- Non-autoregressive: bootstrap multi-step forecast errors at each time point as a whole.

Step 5: Calculate quantiles of the sets of simulated future values at each horizon.

Step 6: Iteratively roll training window forward by one observation; repeat steps 2–5 until prediction intervals have been constructed for the entire test set.

Forecasting heat exposure-related daily mortality



Forecasting heat exposure-related daily mortality


Data

- **Response: Daily deaths in Summer**
– 1990 to 2014 – Montreal, Canada
- **Index Variables:**
 - ▶ Death lags
 - ▶ Max temperature lags
 - ▶ Min temperature lags
 - ▶ Vapor pressure lags
- **Nonlinear:** DOS (day of the season), Year


Data split

- ***tr*** - training window : 1748
- ***tc*** - calibration window : 300
- ***H*** - forecast horizon : 92

Mean coverage of 95% prediction intervals



Mean width of 95% prediction intervals



Summary

- **Block bootstrap** prediction intervals were **narrow** and resulted in **under-coverage**

Summary

- **Block bootstrap** prediction intervals were **narrow** and resulted in **under-coverage**
- **Conformal prediction** methods **generally performed well** across models and forecast horizons

Summary

- **Block bootstrap** prediction intervals were **narrow** and resulted in **under-coverage**
- **Conformal prediction** methods **generally performed well** across models and forecast horizons
- **Conformal bootstrap** method **performed well in shorter horizons**; performance decreased in longer horizons
 - ▶ Using only 1-step errors
 - ▶ Predictor truncation when simulating sample paths


Summary

- **Block bootstrap** prediction intervals were **narrow** and resulted in **under-coverage**
- **Conformal prediction** methods **generally performed well** across models and forecast horizons
- **Conformal bootstrap** method **performed well in shorter horizons**; performance decreased in longer horizons
 - ▶ Using only 1-step errors
 - ▶ Predictor truncation when simulating sample paths
- **Coverage performance** is influenced by: **data characteristics**, **forecasting model**, and **forecast horizon**

Open-source implementation

- Block bootstrap
 - **bb_cvforecast()**
- Conformal bootstrap
 - **cb_cvforecast()**

github.com/nuwani-palihawadana/smimodel



Find me :

-  nuwanipalihawadana.netlify.app
-  [in/nuwani-palihawadana](https://in.linkedin.com/company/nuwani-palihawadana)
-  [@nuwani-palihawadana](https://github.com/nuwani-palihawadana)
-  nuwani.kodikarapalihawadana@monash.edu

References

- Barber, RF, EJ Candès, A Ramdas & RJ Tibshirani (2023). Conformal Prediction Beyond Exchangeability. *The Annals of Statistics* 51(2), 816–845.
- Gibbs, I & E Candès (2021). Adaptive Conformal Inference Under Distribution Shift. *Advances in Neural Information Processing Systems*, 1660–1672.
- Vovk, V, A Gammerman & G Shafer (2005). *Algorithmic Learning in a Random World*. New York, NY: Springer.
- Wang, X & RJ Hyndman (2024). Online Conformal Inference for Multi-Step Time Series Forecasting. *arXiv preprint arXiv:2410.13115*.